

7BDIN006W

Big Data Theory and Practice

Lecture 4

Big Data and Relational Model. SQL: Sorting Data, Subqueries,
Joins

UNIVERSITY OF
WESTMINSTER 

Sorting Data

Using the ORDER BY Clause

- Sort retrieved rows with the ORDER BY clause:
 - ASC: ascending order, default
 - DESC: descending order
- The ORDER BY clause comes last in the SELECT statement:

```
SELECT    ename, job, deptno, hiredate
FROM      emp
ORDER BY  hiredate ;
```

Sorting

- **Sorting in descending order:**

```
SELECT  ename, job, deptno, hiredate  
FROM    emp  
ORDER BY hiredate DESC ;
```

1

- **Sorting by column alias:**

```
SELECT empno, ename, sal*12 annsal  
FROM    emp  
ORDER BY annsal ;
```

2

- **Sorting by multiple columns:**

```
SELECT ename, deptno, sal  
FROM    emp  
ORDER BY deptno, sal DESC;
```

3

Using Subqueries to Solve Queries

Using a Subquery to Solve a Problem

Who has a salary greater than Allen's?

Main query:



Which employees have salaries greater than Allen's salary?

Subquery:



What is Allen's salary?



Subquery Syntax

```
SELECT    select_list
FROM      table
WHERE     expr operator
          (SELECT      select_list
           FROM        table);
```

- The subquery (inner query) executes once before the main query (outer query).
- The result of the subquery is used by the main query.

Using a Subquery

```
SELECT ename, sal  
FROM emp  
WHERE sal >
```

Allen's Salary


```
(SELECT sal  
FROM emp  
WHERE ename = 'ALLEN') ;
```


Guidelines for Using Subqueries

- **Enclose subqueries in parentheses.**
- **Place subqueries on the right side of the comparison condition.**
- **The ORDER BY clause in the subquery is not needed unless you are performing Top-N analysis.**
- **Use single-row operators with single-row subqueries, and use multiple-row operators with multiple-row subqueries.**

Types of Subqueries

- **Single-row subquery**



- **Multiple-row subquery**



Single-Row Subqueries

- Return only one row
- Use single-row comparison operators

Operator	Meaning
=	Equal to
>	Greater than
>=	Greater than or equal to
<	Less than
<=	Less than or equal to
<>	Not equal to

Executing Single-Row Subqueries

```
SELECT ename, job, sal
FROM emp
WHERE job = SALESMAN
AND salary > 1300
      (SELECT sal
       FROM emp
       WHERE empno = 7934) ;
```

The diagram illustrates the execution of a SQL query. The main query is: `SELECT ename, job, sal FROM emp WHERE job = SALESMAN AND salary > 1300`. The subquery `(SELECT job FROM emp WHERE ename = 'ALLEN')` is shown in a red box, and its result `SALESMAN` is highlighted in red. A red arrow points from this result to the `job =` condition in the main query. Another subquery `(SELECT sal FROM emp WHERE empno = 7934)` is shown in a red box, and its result `1300` is highlighted in red. A red arrow points from this result to the `salary >` condition in the main query.

Will This Statement Return Rows?

```
SELECT  ename, job
FROM    emp
WHERE   job =
        (SELECT job
         FROM    emp
         WHERE   ename = 'BING') ;
```

no rows selected

Subquery returns no values.

What Is Wrong with This Statement?

```
SELECT empno, ename  
FROM emp  
WHERE job =
```

```
(SELECT job  
FROM emp  
WHERE deptno = 20)
```

```
ERROR at line 4:  
ORA-01427: single-row subquery returns more than  
one row
```

Single-row operator with multiple-row subquery

Multiple-Row Subqueries

- Return more than one row
- Use multiple-row comparison operators

Operator	Meaning
IN	Equal to any member in the list
∃ ANY ∃ SOME	Compare value to each value returned by the subquery
∃ ALL	Compare value to every value returned by the subquery
where ∃ can be any of the 6 comparison operators, i.e. < <= = >= > and <> or !=	

Using the IN or =ANY Operator in Multiple-Row Subqueries

```
SELECT empno, ename
FROM emp
WHERE job =
      (SELECT job
       FROM emp
       WHERE deptno = 20)
```

Above returns an error as the subquery returns multiple rows

```
SELECT empno, ename
FROM emp
WHERE job IN
      (SELECT job
       FROM emp
       WHERE deptno = 20)
```

```
SELECT empno, ename
FROM emp
WHERE job = ANY
      (SELECT job
       FROM emp
       WHERE deptno = 20)
```


Using the ANY Operator in Multiple-Row Subqueries

```
SELECT empno, ename, job, sal
FROM emp
WHERE sal < ANY (SELECT sal
                  FROM emp
                  WHERE job = 'SALESMAN')
AND job <> 'SALESMAN';
```

1250, 1600, 1500

Using the ALL Operator in Multiple-Row Subqueries

```
SELECT empno, ename, job, sal
FROM emp
WHERE sal < ALL (SELECT sal
                  FROM emp
                  WHERE job = 'SALESMAN')
AND job <> 'SALESMAN';
```

1250, 1600, 1500

Executing a Multi-Column Single-Row Subqueries

```
SELECT ename, job, sal
FROM emp
WHERE (job, sal) = (ANALYST, 3000)
                (SELECT job, sal
                 FROM emp
                 WHERE ename = 'SCOTT');
```

The diagram illustrates a multi-column single-row subquery. The main query selects employee details from the 'emp' table, filtering by job and salary. The subquery, enclosed in a red box, selects the job and salary for the employee named 'SCOTT'. A red arrow points from the subquery's result to the comparison operator in the WHERE clause, indicating that the subquery's output is used to filter the main query's results.

Executing a Multi-Column Single-Row Subqueries

```
SELECT ename, job, sal
FROM emp
WHERE (job, sal) =ANY
      (SELECT job, sal
       FROM emp
       WHERE deptno = 30) ;
```

only the =ANY and the IN can be used

```
SELECT ename, job, sal
FROM emp
WHERE (job, sal) <= ANY
      (SELECT job, sal
       FROM emp
       WHERE deptno = 30) ;
```

Null Values in a Subquery

```
SELECT  ename
FROM    emp
WHERE   empno NOT IN
        (SELECT mgr
         FROM    emp
         WHERE   job = 'PRESIDENT') ;
```

no rows selected

Nested Subqueries

```
SELECT empno, ename, job, sal
FROM emp
WHERE sal < ALL
```

```
(SELECT sal
  FROM emp
 WHERE job = 'SALESMAN'
      AND deptno = ANY
```

```
(SELECT deptno
  FROM dept
 WHERE loc = 'CHICAGO'))
```

```
AND job <> 'SALESMAN';
```

Displaying Data from Multiple Tables

Types of Joins

Joins that are compliant with the SQL:1999 standard include the following:

- **Cross joins**
- **Natural joins**
- **USING clause**
- **Full (or two-sided) outer joins**
- **Arbitrary join conditions for outer joins**

Joining Tables Using SQL:1999 Syntax

Use a join to query data from more than one table:

```
SELECT    table1.column, table2.column
FROM      table1 [, table2] |
          [CROSS JOIN table2] |
          [JOIN table2
            ON (table1.column_name = table2.column_name)] |
          [JOIN table2
            ON (condition_involving table1 and table2)] |
          [NATURAL JOIN table2] |
          [JOIN table2 USING (column_name)] |
          [LEFT|RIGHT|FULL OUTER JOIN table2
            ON (table1.column_name = table2.column_name)] ;
```

Cartesian Products

- **A Cartesian product is formed when:**
 - **A join condition is omitted**
 - **A join condition is invalid**
 - **All rows in the first table are joined to all rows in the second table**
- **To avoid a Cartesian product, always include a valid join condition.**

Creating Cross Joins

- The CROSS JOIN clause produces the cross-product of two tables.
- This is also called a Cartesian product between the two tables.

```
SELECT *  
FROM   emp  
       CROSS JOIN dept ;
```

```
SELECT emp.*, dname  
FROM   emp  
       CROSS JOIN dept  
WHERE  emp.deptno = dept.deptno;
```

```
SELECT emp.*, dname  
FROM   emp  
       CROSS JOIN dept  
WHERE  emp.deptno = dept.deptno  
and    loc = 'CHICAGO';
```

Qualifying Ambiguous Column Names

- **Use table prefixes to qualify column names that are in multiple tables.**
- **Use table prefixes to improve performance.**
- **Use column aliases to distinguish columns that have identical names but reside in different tables.**
- **Do not use aliases on columns that are identified in the `USING` clause and listed elsewhere in the SQL statement.**

Using Table Aliases

- Use table aliases to simplify queries.
- Use table aliases to improve performance.

```
SELECT e.empno, e.ename, d.loc, d.deptno
FROM   emp e
       CROSS JOIN dept d
WHERE  e.deptno = d.deptno
       and loc = 'CHICAGO';
```

Creating Joins with the ON Clause

- The join condition for the natural join is basically an equijoin of all columns with the same name.
- Use the ON clause to specify arbitrary conditions or specify columns to join.
- The join condition is separated from other search conditions.
- The ON clause makes code easy to understand.

Retrieving Records with the ON Clause

```
SELECT empno, ename, e.deptno, d.deptno, loc  
FROM   emp e  
       JOIN dept d  
         ON e.deptno = d.deptno);
```

Self-Joins Using the ON Clause

You want to find the details of employees along with those of their managers

EMP (WORKER)

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO
7369	SMITH	CLERK	7902	17-Dec-1980	800	-	20
7499	ALLEN	SALESMAN	7698	20-Feb-1981	1600	800	30
7521	WARD	SALESMAN	7698	22-Feb-1981	1250	500	30
7566	JONES	MANAGER	7839	02-Apr-1981			
7654	MARTIN	SALESMAN	7698	20-Sep-1981			

EMP (MANAGER)

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO
7369	SMITH	CLERK	7902	17-Dec-1980	800	-	20
7499	ALLEN	SALESMAN	7698	20-Feb-1981	1600	800	30
7521	WARD	SALESMAN	7698	22-Feb-1981	1250	500	30
7566	JONES	MANAGER	7839	02-Apr-1981	2975	-	20
7654	MARTIN	SALESMAN	7698	20-Sep-1981	1250	1400	30

MGR in the *WORKER* table is equal to EMPNO in the *MANAGER* table.

Self-Joins Using the ON Clause

```
SELECT e.empno emp_empno, e.ename emp_ename, e.mgr emp_mgrno,  
       m.empno man_empno, m.ename man_ename  
FROM   emp e  
       JOIN emp m  
       ON (e.mgr = m.empno);
```

Applying Additional Conditions to a Join

```
SELECT e.empno, e.ename, d.loc, d.deptno
FROM   emp e
       JOIN dept d
       ON (e.deptno = d.deptno)
       and loc = 'CHICAGO';
```

```
SELECT e.empno, e.ename, d.loc, d.deptno
FROM   emp e
       JOIN dept d
       ON (e.deptno = d.deptno)
WHERE  loc = 'CHICAGO';
```

```
SELECT e.empno, e.ename, d.loc, d.deptno
FROM   emp e
       JOIN dept d
       ON (e.deptno = d.deptno
          and 'CHICAGO');
```

Retrieving Records with Nonequijoins

```
SELECT e.ename, sal, grade  
FROM   emp e  
       JOIN salgrade  
         ON sal BETWEEN  
            losal AND hisal;
```

Creating Three-Way Joins with the ON Clause

```
SELECT empno, ename, e.deptno, dname, grade
FROM   emp e
       JOIN dept d
         ON d.deptno = e.deptno
       JOIN salgrade
         ON (sal BETWEEN
            losal AND hisal);
```